

Spencer Sullivan

(512) 922-2434 | sas@fsullivans.net | San Diego Area | spencer-sullivan.com

SUMMARY

Software / Data / Machine Learning Engineer with 7+ years of experience in building business-critical applications that rely heavily on data and modeling. Currently helping ScaleFactor to automate accounting for small businesses.

IDEAL JOB

I am looking for opportunities in the San Diego area (or remote) where I can significantly impact the business, work closely with researchers / data scientists to develop, productionize, and scale models, and continually learn and grow as an engineer.

KEY SKILLS

- **Platforms:** Amazon Web Services (AWS), Google Cloud, On-Premise Cluster
- **Programming Languages:** Ruby (Rails), Python, Go, SQL, HTML, JS (Ember)
- **Databases:** Snowflake, Presto, Athena, Redshift, Postgres, Redis, Dynamo
- **Data Tools & Frameworks:** Airflow, Luigi, Glue, Hive, Spark, Kafka / Kinesis
- **Other:** Docker, Git, Kubernetes, DataDog, Tableau, Lambda, SQS

CAREER

Senior Software Engineer	Remote
<i>ScaleFactor</i>	2019 – Present
Automating small-business accounting with full-stack web development and ML pipelining.	

Senior Software Engineer	Austin, TX
<i>Spiceworks</i>	2017 – 2019
Architected and developed Petabyte-scale data pipelines feeding ML models.	

Software Engineer	Austin, TX
<i>Spiceworks</i>	2014 – 2017
Architected and developed Petabyte-scale data pipelines feeding ML models.	

Software Development Intern	Austin, TX
<i>Spiceworks</i>	2010 – 2014
Architected and developed Petabyte-scale data pipelines feeding ML models.	

EDUCATION

University of Texas at Austin

Austin, TX

Bachelors of Science in Computer Science

2011 - 2014

EXPERIENCE

ML for Bookkeeping *An ML workflow for automatically classifying bank transactions*

- Reduced need for humans to classify transactions manually
- Refined frontend, backend, and modeling code for production use
- Analyzed model performance

Accounting Report Card *Self-service health dashboard for business accounting books*

- Integrated with 3rd parties to collect and parse accounting reports
- Ran analysis to determine healthy accounting patterns
- Created portal for calculating and viewing insights

Real Time Bidder *Set of services for intelligently bidding on ad inventory*

- Wrote services that respond in under 10ms (including network latency)
- Horizontally scaled web server up to 1 billion requests per day
- Designed micro-service architecture comprising 10+ services
- Productionized machine learning model to compute bids
- Helped scale product from \$0 to ~\$10M in annual revenue
- Rigorously tested and wrote monitoring

Business Identity Graph *Database of companies who visited Spiceworks*

- Generated graph from millions of data points from several sources
- Designed data model where graph was iteratively updated in-place
- Resolved business attributes from a connected component of the graph

Account Intelligence Dashboard *showing businesses and their interests*

- Wrote data pipeline to calculate purchase intent of businesses
- Integrated Tensorflow model from data science team into pipeline
- Architected pipeline to leverage business identity graph

Pixel Tracking *Web-tracking service integrated with 100's of companies*

- Led rewrite of service
- Integrated open-source browser fingerprinting

Segmentation Service *Puts people into segments based on their actions*

- Wrote jobs to send billions of events into system every month
- Integrated with 3rd parties using cookie matching
- Debugged / resolved problems around user-identity mappings
- Collaborated with Data Science to incorporate intent modeling