Spencer Sullivan

(512) 922-2434 │ sas@fsullivans.net │ Austin, TX │ spencer-sullivan.com

SUMMARY	Software / Data / Machine Learning Engineer with 7+ years of experience in building business-critical applications that rely heavily on data and modeling. Currently helping ScaleFactor to automate accounting for small businesses.

IDEAL JOB

I am looking for opportunities in the San Diego area (or remote) where I can significantly impact the business, work closely with researchers / data scientists to develop, productionize, and scale models, and continually learn and grow as an engineer.

KEY SKILLS

Platforms: Amazon Web Services (AWS), Google Cloud, On-Premise Cluster
Programming Languages: Ruby (Rails), Python, Go, SQL, HTML, JS (Ember)
Databases: Snowflake, Presto, Athena, Redshift, Postgres, Redis, Dynamo
Data Tools & Frameworks: Airflow, Luigi, Glue, Hive, Spark, Kafka / Kinesis
Other: Docker, Git, Kubernetes, DataDog, Tableau, Lambda, SQS

CAREER

Senior Software Engineer   2019 - Present
ScaleFactor, Remote
Automating small-business accounting with full-stack web development and ML pipelining.

Senior Software Engineer   2017 - 2019
Spiceworks, Austin
Architected and developed Petabyte-scale data pipelines feeding ML models.

Software Engineer   2014 - 2017
Spiceworks, Austin
Wrote business-critical, high-volume, low-latency systems, mostly in Golang.

Software Development Intern   2010 - 2014
Spiceworks, Austin
Learned how to work in agile teams and translate business requirements to code.

EDUCATION

BS in Computer Science, *University of Texas at Austin*   2011 - 2014

Experience	**ML for Bookkeeping** *An ML workflow for automatically classifying bank transactions*

ML for Bookkeeping *An ML workflow for automatically classifying bank transactions*
- Reduced need for humans to classify transactions manually
- Refined frontend, backend, and modeling code for production use
- Analyzed model performance

Accounting Report Card *Self-service health summary dashboard of business accounting books*
- Integrated with 3rd parties to collect and parse accounting reports
- Ran analysis to determine healthy accounting patterns
- Created portal for calculating and viewing insights

Real Time Bidder *Set of services for intelligently bidding on ad inventory*
- Wrote services that respond in under 10ms (including network latency)
- Horizontally scaled web server up to 1 billion requests per day
- Designed micro-service architecture comprising 10+ services
- Productionized machine learning model to compute bids
- Helped scale product from $0 to ~$10M in annual revenue
- Rigorously tested and wrote monitoring

Business Identity Graph *Database of companies who visited Spiceworks*
- Generated graph from millions of data points from several sources
- Designed data model where graph was iteratively updated in-place
- Resolved business attributes from a connected component of the graph

Account Intelligence *Dashboard showing businesses and their interests*
- Wrote data pipeline to calculate purchase intent of businesses
- Integrated Tensorflow model from data science team into pipeline
- Architected pipeline to leverage business identity graph

Pixel Tracking *Web-tracking service integrated with 100's of companies*
- Led rewrite of service
- Integrated open-source browser fingerprinting

Segmentation service *Puts people into segments based on their actions*
- Wrote jobs to send billions of events into system every month
- Integrated with 3rd parties using cookie matching
- Debugged / resolved problems around user-identity mappings
- Collaborated with Data Science to incorporate intent modeling